









# Compute Service

At the end of this module, you should be able to:

- Identify how Amazon Elastic Compute Cloud (Amazon EC2) can be used.
- Explain the value of using Amazon Machine Images (AMIs) to accelerate the creation and repeatability of infrastructure
- Differentiate between the EC2 instance types
- Recognize storage solutions for Amazon EC2
- Describe EC2 pricing options
- Describe AWS Lambda service
- Describe Elastic Beanstalk service

# Adding compute with Amazon EC2

# AWS runtime compute choices

| Virtual Machines (VMs)   | Containers   | Platform as a Service (PaaS)   | Serverless   | Specialized Solutions   |
|--|--|--|--|---|
| <br>Amazon Elastic Compute Cloud (Amazon EC2) | <br>Amazon Elastic Container Service (Amazon ECS) | <br>AWS Elastic Beanstalk | <br>AWS Lambda  | <br>AWS Outposts |
| <br>Amazon Lightsail                          |  |  | <br>AWS Fargate | <br>AWS Batch    |


Higher infrastructure control and customization

Faster application deployment

Fully managed services

Different compute services are available to meet the needs of different use cases.

This module will discuss Amazon EC2.



Amazon Elastic  
Compute Cloud  
(Amazon EC2)


**Amazon EC2** provides resizable compute capacity in the cloud.

- Provides virtual machines (servers)
- Provisions servers in minutes
- Can automatically scale capacity up or down as needed
- Enables you to pay only for the capacity that you use

# EC2 instances

An EC2 instance is a **virtual machine** that runs on a physical host.

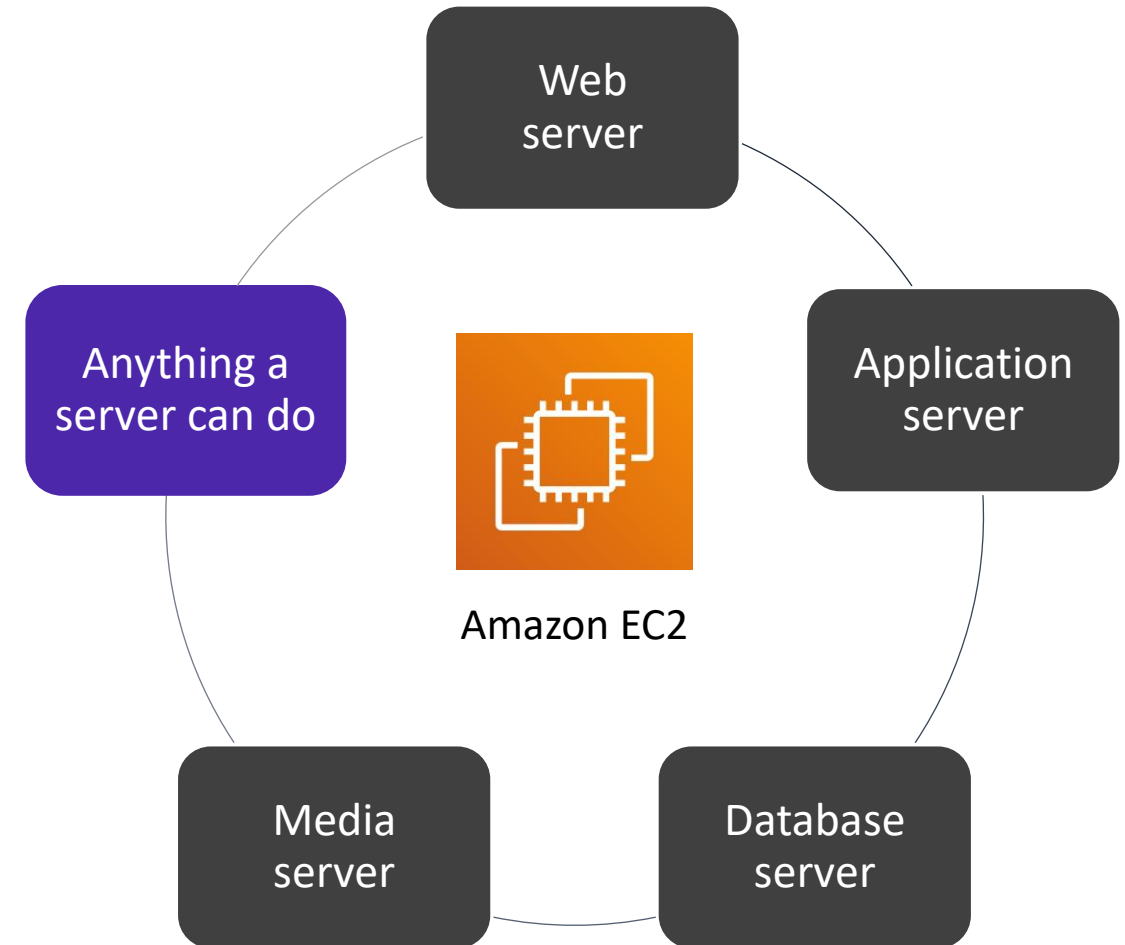
- You can choose different configurations of CPU and memory capacity
- Supports different storage options
  - Instance store
  - Amazon Elastic Block Store (Amazon EBS)
- Provides network connectivity



# Amazon EC2 use cases

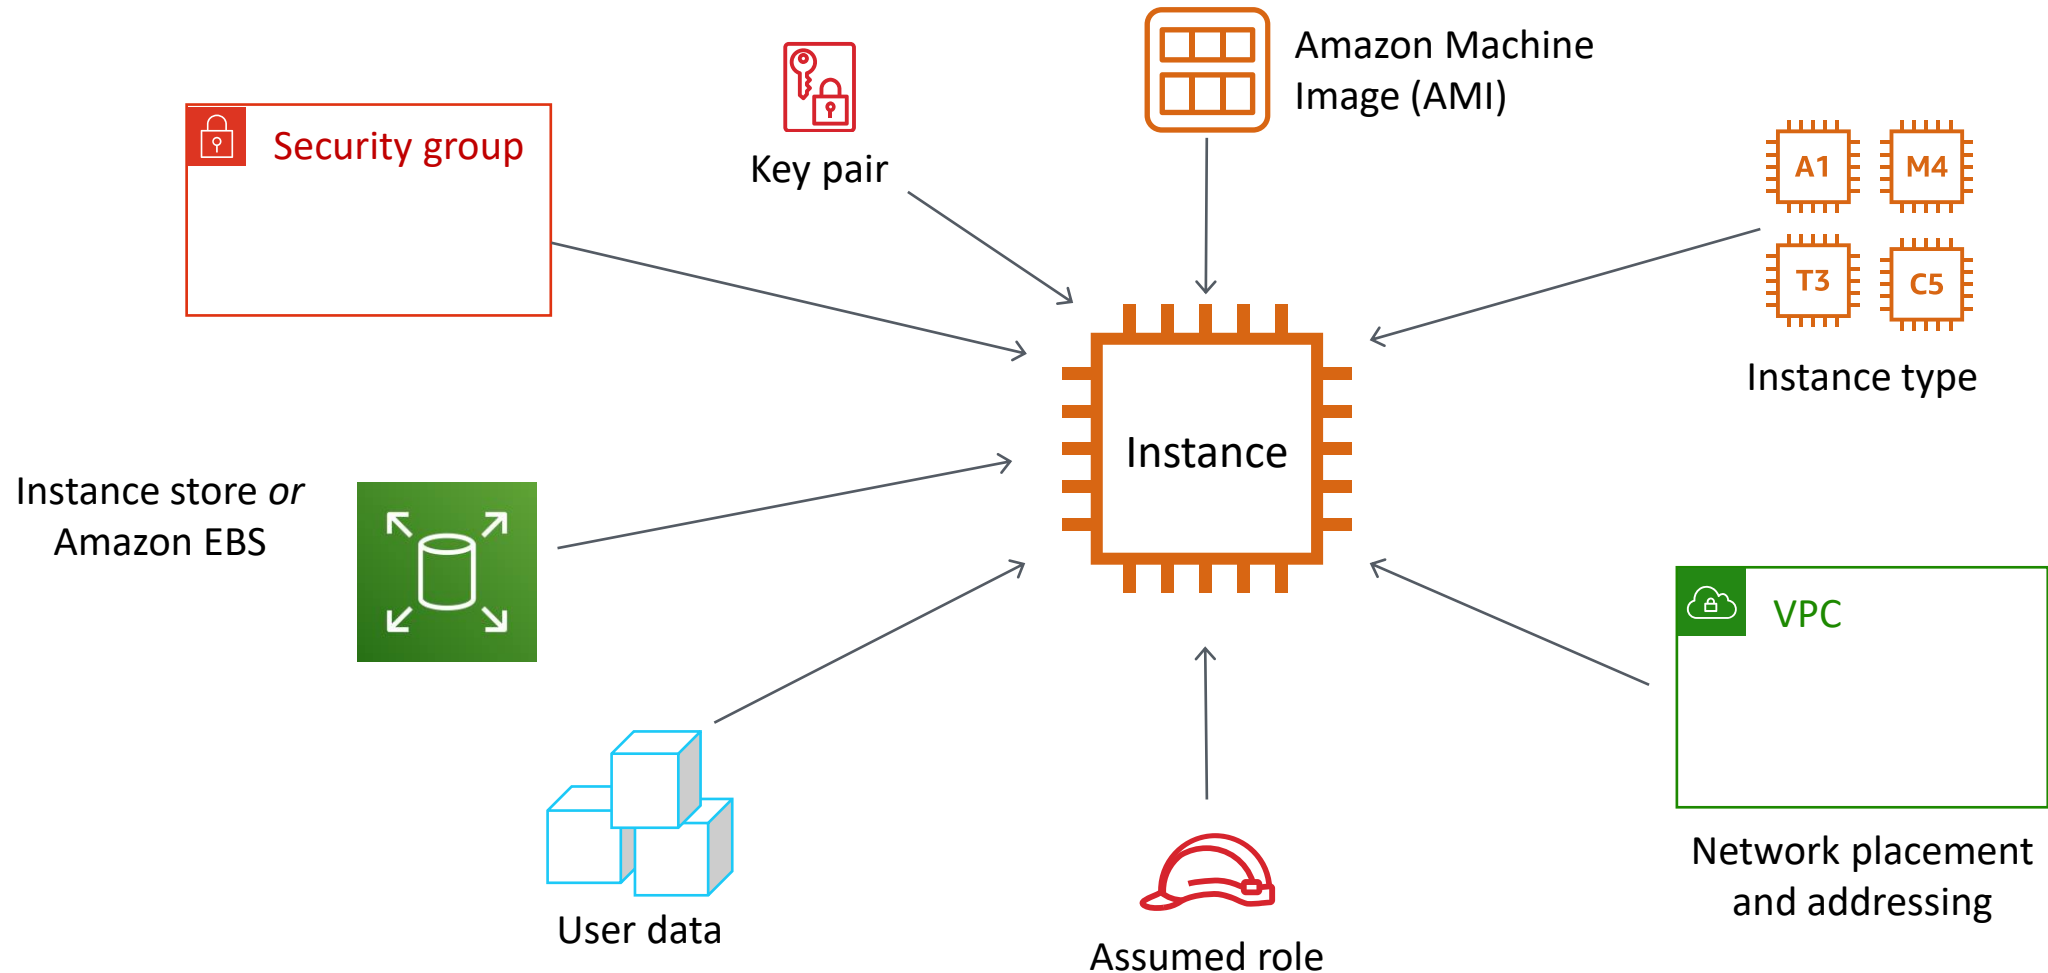
Use Amazon EC2 when you need:

- Complete control of your computing resources, including *operating system* and *processor type*
- Options for optimizing your compute costs –
  - *On-Demand Instances*, *Reserved Instances*, and *Spot Instances*
  - *Savings Plans*
- Ability to run any type of workload, for example –
  - Simple websites
  - Enterprise applications
  - High performance computing (HPC) applications



# Provisioning an EC2 instance

## Essential instance launch configuration parameters





# key takeaways



- Amazon EC2 enables you to run Microsoft Windows and Linux **virtual machines** in the cloud.
- You can use an EC2 instance when you need **complete control of your computing resources** and want to **run any type of workload**.
- When you launch an EC2 instance, you must choose an **AMI** and an **instance type**. Launching an instance involves specifying configuration parameters, including **network, security, storage, and user data settings**.

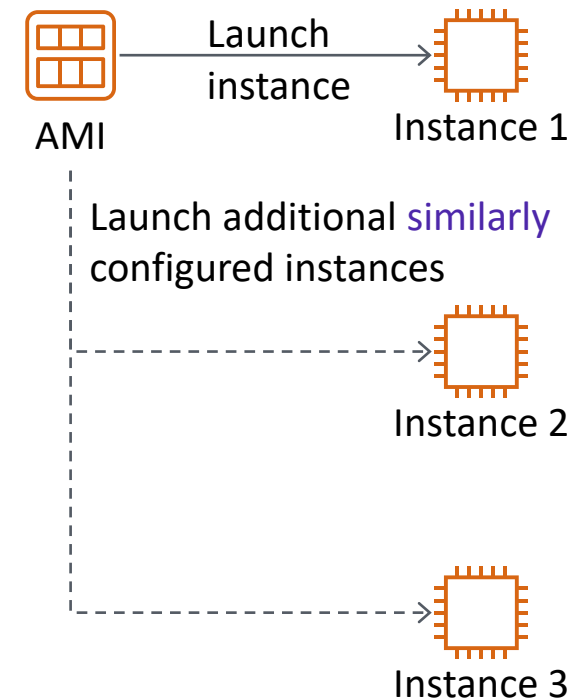
# Choosing an AMI to launch an EC2 instance

# Amazon Machine Image (AMI)

An **AMI** provides the information that is needed to launch an instance, including:

- A **template** for the root volume
  - Contains the guest operating system (OS) and perhaps other installed software
- **Launch permissions**
  - Control which AWS accounts can access the AMI
- **Block device mappings**
  - Specifies any storage volumes to attach to the instance

Create multiple instances from the same AMI






- **Repeatability**
  - An AMI can be used repeatedly to launch instances with efficiency and precision
- **Reusability**
  - Instances launched from the same AMI are identically configured
- **Recoverability**
  - You can create an AMI from a configured instance as a restorable backup
  - You can replace a failed instance by launching a new instance from the same AMI

# Choosing an AMI

Choose an AMI based on:

- Region
- Operating system
  - Microsoft Windows or Linux
- Storage type of the root device
- Architecture
- Virtualization type



AMI sources:

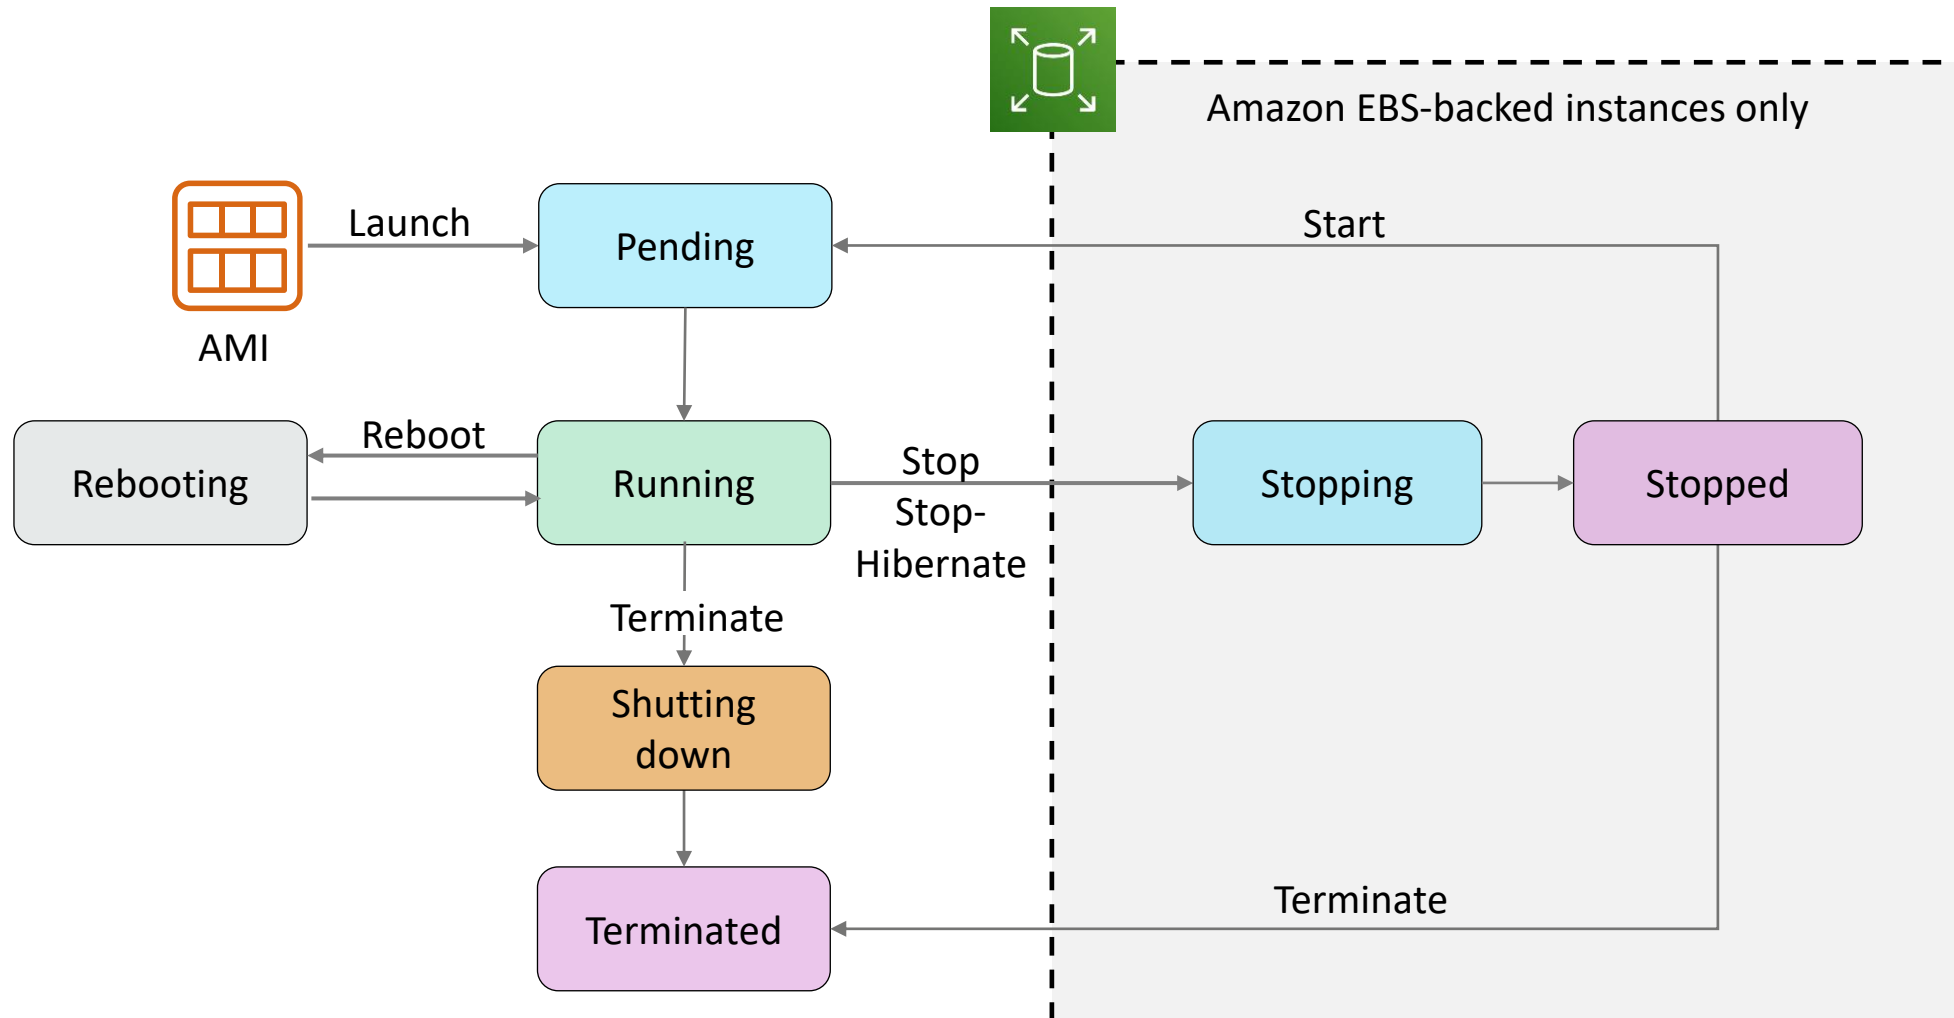
- [Quick Start](#) – *Linux and Microsoft Windows AMIs that are provided by AWS.*
- [My AMIs](#) – *Any AMIs that you create.*
- [AWS Marketplace](#) – *Pre-configured templates from third parties.*
- [Community AMIs](#) – *AMIs shared by others. Use at your own risk.*



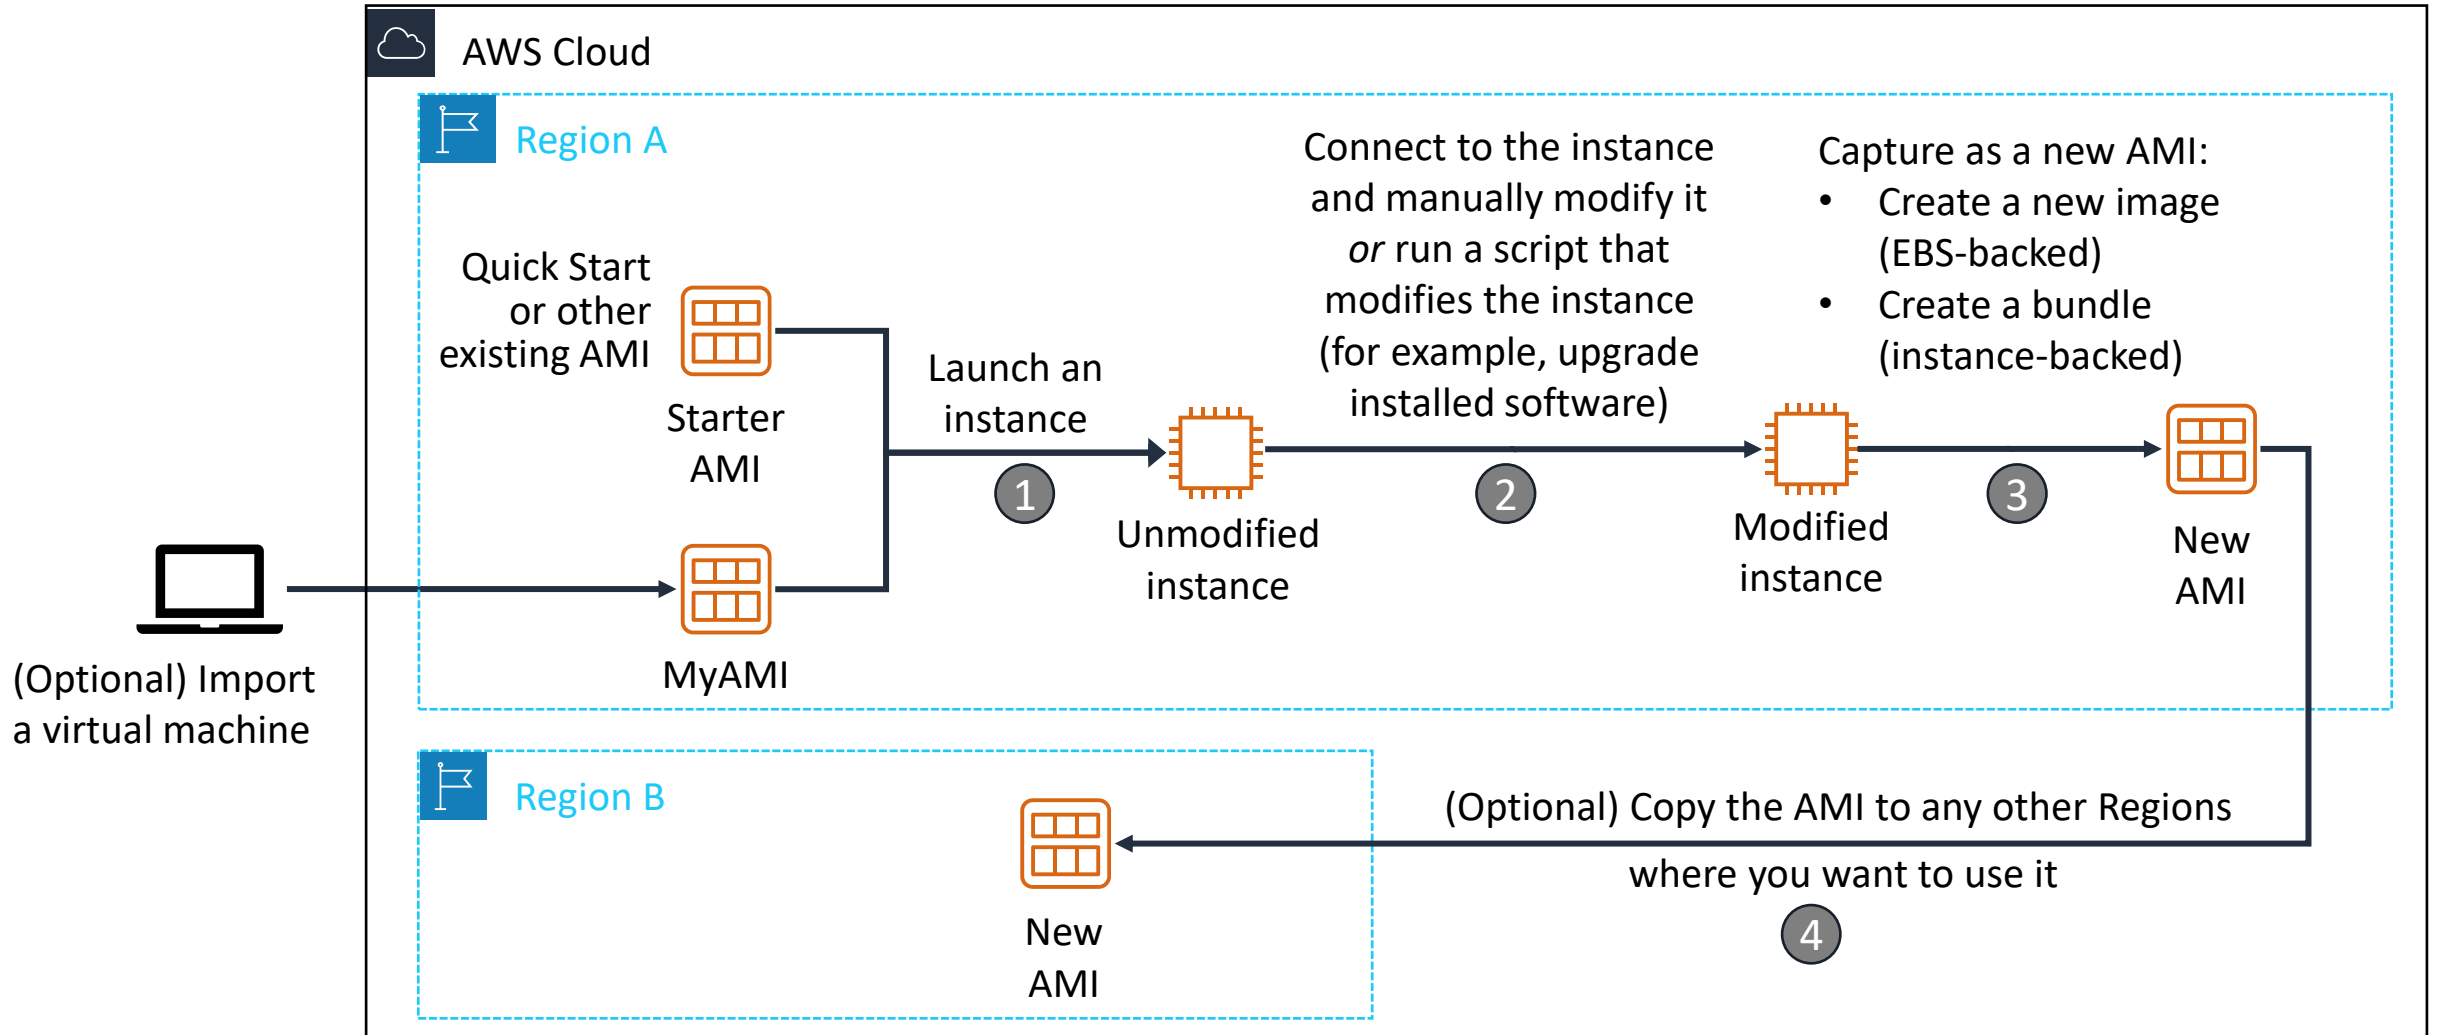
# Instance store-backed versus Amazon EBS-backed AMI

| Characteristic                      | Amazon EBS-Backed Instance  | Instance Store-Backed Instance                                       |
|-------------------------------------|---|--|
| Boot time for the instance          | Boots faster  | Takes longer to boot   |
| Maximum size of root device         | 16 TiB  | 10 GiB   |
| Ability to stop the instance        | Can stop the instance   | Can't stop the instance, only reboot or terminate it                 |
| Ability to change the instance type | Can change the instance type by stopping instance   | Can't change the instance type because the instance can't be stopped |
| Instance charges                    | You are charged for instance usage, EBS volume usage, and storing your AMI as an EBS snapshot | You are charged for instance usage and storing your AMI in Amazon S3 |

# Amazon EC2 instance lifecycle



# Creating a new AMI







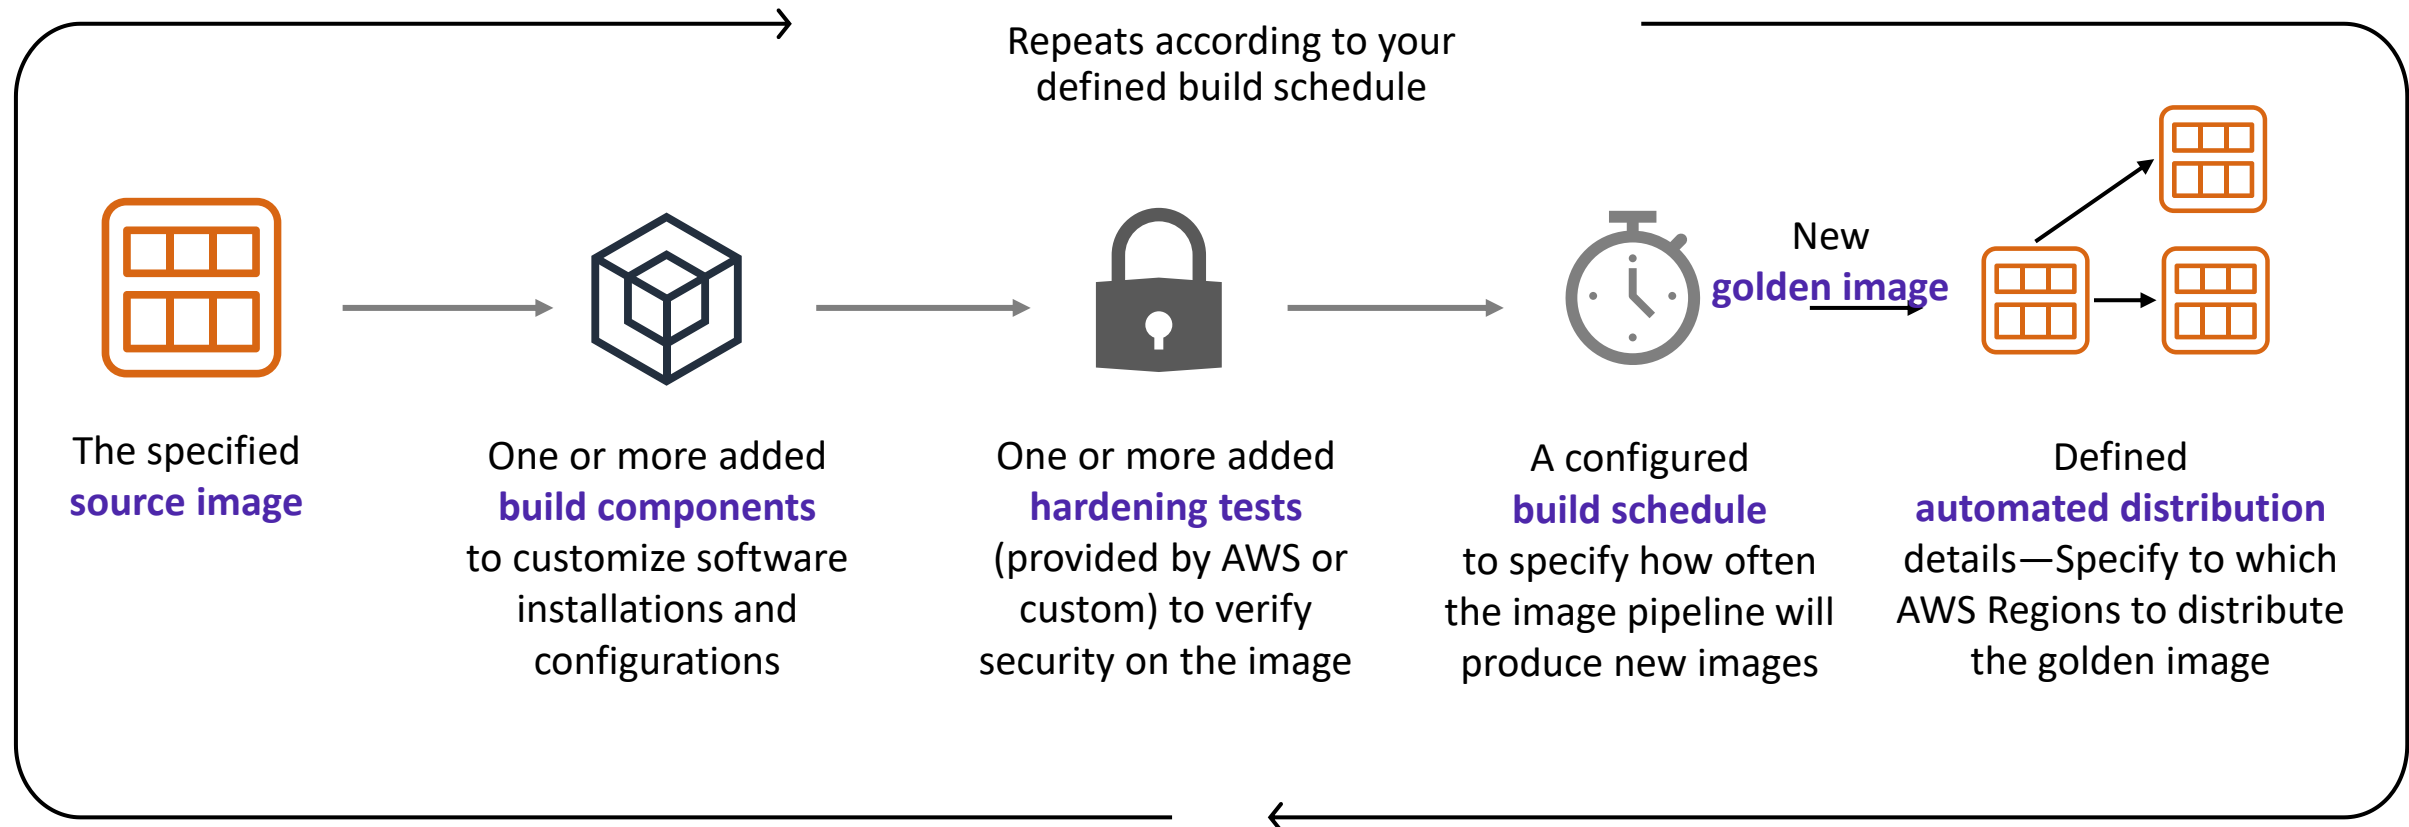
EC2 Image Builder

EC2 Image Builder automates the **creation, management, and deployment** of up-to-date and compliant **golden VM images**.

- Provides a graphical interface to create image-building pipelines
- Creates and maintains **Amazon EC2 AMIs** and on-premises VM images
- Produces secure, validated, and up-to-date images
- Enforces version control

# How EC2 Image Builder works

## An EC2 Image Builder **image pipeline**



# key takeaways

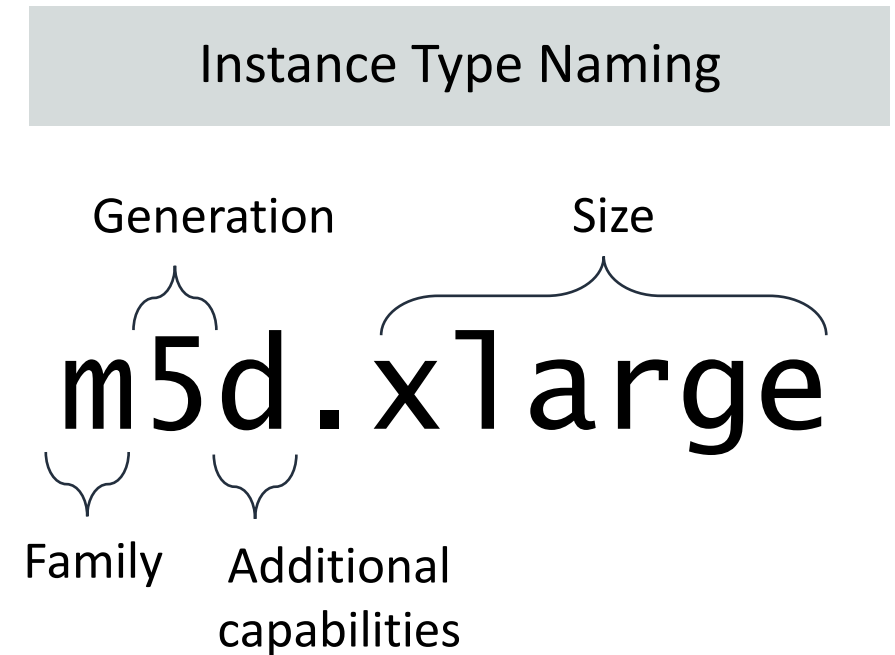



- An **AMI** provides the information that is needed to launch an EC2 instance
- For best performance, use an AMI with **HVM virtualization type**
- Only an instance launched from an Amazon EBS-backed AMI **can be stopped and started**
- An AMI is available in a **Region**

# Selecting an EC2 instance type

# EC2 instance type

An **EC2 instance type** defines a configuration of CPU, memory, storage, and network performance characteristics that provides a given level of compute performance.




# Suitability of instance types for workloads (1 of 2)

## General purpose instance types

- Web or application servers
- Enterprise applications
- Gaming servers
- Analytics applications
- Development or test environments


Example instance types:



## Compute optimized instance types

- Batch processing
- Distributed analytics
- High performance computing (HPC)
- Ad server engines
- Multiplayer gaming
- Video encoding

Example instance types:




# Suitability of instance types for workloads (2 of 2)

## Memory optimized instance types

- In-memory caches
- High-performance databases
- Big data analytics

Example instance types:



## Accelerated computing instance types

- Machine learning, artificial intelligence (AI)
- HPC
- Graphics

Example instance types:



## Storage optimized instance types

- High-performance databases<sup>1</sup>
- Real-time analytics<sup>1</sup>
- Transactional workloads<sup>1</sup>
- NoSQL databases<sup>1</sup>
- Big data<sup>2</sup>
- Data warehouse<sup>2</sup>
- Log processing<sup>2</sup>

<sup>1</sup>High I/O example instance type:



<sup>2</sup>Dense Storage example instance types:



# Choosing an instance type

- Choose the instance type that meets –
  - The **performance needs** of your application
  - Your **cost requirements**
- When you create a *new* instance –
  - In the EC2 console, use the **Instance Types** page to filter by characteristics that you choose
  - Recommendation: The latest generation in an instance family typically has a better price-to-performance ratio
- If you have an *already existing* instance –
  - You can get recommendations for optimizing the instance type by using the [AWS Compute Optimizer](#)
  - You can evaluate recommendations and modify the instance accordingly

With over 270 available instances types, how do you choose the correct one?





# AWS Compute Optimizer

## AWS Compute Optimizer

- Recommends *optimal* instance type, instance size, and Auto Scaling group configuration
- Analyzes workload patterns and makes recommendations
- Classifies instance findings as Under-provisioned, Over-provisioned, Optimized, or None

AWS Compute Optimizer > Dashboard > Recommendations for EC2 instances

**Recommendations for EC2 instances (8)** [Info](#)

Recommendations for modifying current resources for better cost and performance.

Filter by one or more Regions  Over-provisioned < 1 > [View detail](#)

Region: US East (N. Virginia) X Clear filters

|                       | Instance ID ▲       | Instance name ▼ | Finding ▼        | Current instance type ▼ | Current On-Demand price ▼ | Recommended instance type ▼ |
|-----------------------|---------------------|-----------------|------------------|-------------------------|---------------------------|-----------------------------|
| <input type="radio"/> | i-0218a45abd8b53658 | -               | Over-provisioned | m5.xlarge               | \$0.192 per hour          | r5.large                    |
| <input type="radio"/> | i-069f6e837890db127 | -               | Over-provisioned | c5.xlarge               | \$0.17 per hour           | t3.large                    |
| <input type="radio"/> | i-07084b94d1bcf391b | -               | Over-provisioned | c5.xlarge               | \$0.17 per hour           | t3.large                    |
| <input type="radio"/> | i-0af9322ff627d7e8f | -               | Over-provisioned | m5.xlarge               | \$0.192 per hour          | r5.large                    |
| <input type="radio"/> | i-0ceb95ed248026d24 | -               | Over-provisioned | m5.xlarge               | \$0.192 per hour          | r5.large                    |
| <input type="radio"/> | i-0f277818dfef522e9 | -               | Over-provisioned | c5.xlarge               | \$0.17 per hour           | t3.large                    |
| <input type="radio"/> | i-0f4f4c06ad8afe81a | -               | Over-provisioned | m5.2xlarge              | \$0.384 per hour          | r5.xlarge                   |
| <input type="radio"/> | i-0fb9323080785de1e | -               | Over-provisioned | c5.xlarge               | \$0.17 per hour           | t3.large                    |

# key takeaways



- An **EC2 instance type** defines a configuration of CPU, memory, storage, and network performance characteristics
- As a recommendation, choose **new generation instance types in a family** because they generally have better price-to-performance ratios
- Use the **Instance Types** page in the Amazon EC2 console and **AWS Compute Optimizer** to find the right instance type for your workload

# Adding storage to an Amazon EC2 instance

# Amazon EC2 storage overview

## Root volume

This volume always contains the guest OS



Instance store



Amazon EBS  
(SSD-backed only)



An EC2 instance will *always* have a **root volume**, and can *optionally* have one or more **data volumes**.

## Data volumes

For data accessed by a single instance



Instance store



Amazon EBS

For data accessible from multiple instances




Amazon Elastic File System  
(Amazon EFS) [Linux]



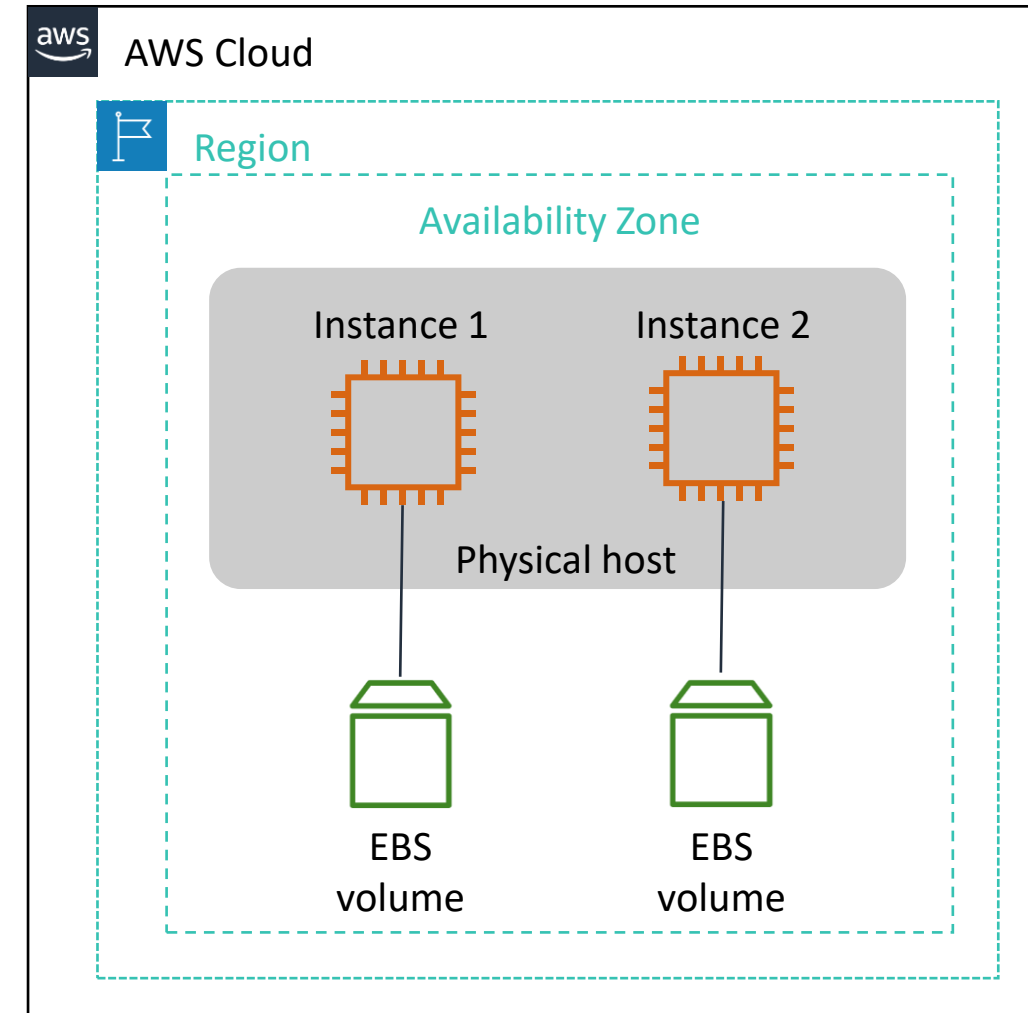
Amazon FSx for Windows  
File Server

# Instance store

- An instance store provides **non-persistent storage** to an instance –
  - The data is stored on the *same physical server* where the instance runs
- Characteristics –
  - Temporary block-level storage
  - Uses HDD or SSD
  - **Instance store data is lost when the instance is *stopped* or *terminated***
- Example use cases –
  - Buffers
  - Cache
  - Scratch data



- Amazon EBS volumes provide **network-attached persistent storage** to an EC2 instance.
- Characteristics –
  - Is persistent block-level storage
  - Can attach to any instance in the same Availability Zone
  - Uses HDD or SSD
  - Can be encrypted
  - Supports snapshots that are persisted to S3
  - Data persists independently from the life of the instance
- Example use cases –
  - Stand-alone database
  - General application data storage



# Amazon EBS SSD-backed volume types

Amazon EBS SSD-backed volumes are suited for use cases where the performance focus is on IOPS.

|             | General Purpose SSD (gp2)   | Provisioned IOPS SSD (io1)  |
|-------------|---|---|
| Description | Balances price and performance for a wide variety of workloads  | <ul style="list-style-type: none"><li>• Highest-performance SSD volume</li><li>• Good for mission-critical, low-latency, or high-throughput workloads</li></ul>   |
| Use Cases   | <ul style="list-style-type: none"><li>• Recommended for most workloads</li><li>• Can be a boot volume</li></ul> | <ul style="list-style-type: none"><li>• Critical business applications that require sustained IOPS performance</li><li>• Large database workloads</li><li>• Transactional workloads</li><li>• It can be a boot volume</li></ul> |

# Amazon EBS HDD-backed volume types

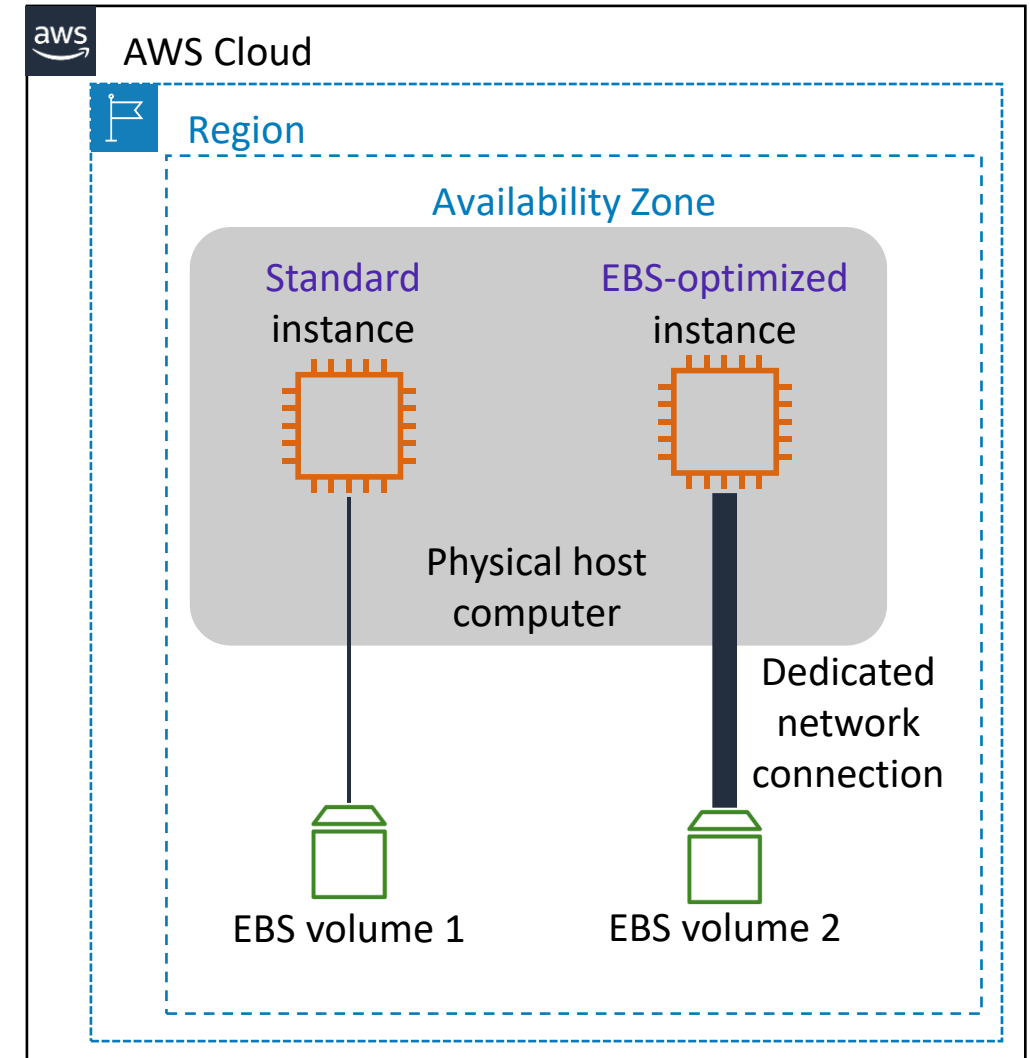
Amazon EBS HDD-backed volumes work well when the focus is on throughput.

|             | Throughput Optimized HDD (st1)   | Cold HDD (sc1)  |
|-------------|--|---|
| Description | <ul style="list-style-type: none"><li>• Low-cost volume type</li><li>• Designed for frequently accessed, throughput-intensive workloads</li></ul>                                | <ul style="list-style-type: none"><li>• Lowest-cost HDD volume</li><li>• Designed for less frequently accessed workloads</li></ul>  |
| Use Cases   | <ul style="list-style-type: none"><li>• Streaming workloads</li><li>• Big data</li><li>• Data warehouses</li><li>• Log processing</li><li>• It cannot be a boot volume</li></ul> | <ul style="list-style-type: none"><li>• Throughput-oriented storage for large volumes of infrequently accessed data</li><li>• Use cases where the lowest storage cost is important</li><li>• It cannot be a boot volume</li></ul> |



# Amazon EBS-optimized instances

- Certain EC2 instance types can be **EBS-optimized**
- Benefits –
  - Provides a **dedicated network connection** to attached EBS volumes
  - Increases I/O performance
  - Additional performance is achieved if using an Amazon EC2 **Nitro System**-based instance type
- Usage –
  - For EBS-optimized instance types, optimization is enabled by default
  - For other instances types that support it, optimization must be manually enabled



# Shared file systems for EC2 instances

What if you have **multiple instances** that must use the **same storage**?

Amazon EBS: Attaches only to one instance



Amazon EBS

Amazon S3: Is an option, but is not ideal



Amazon S3

Amazon EFS *and* Amazon FSx for Windows File Server: Both satisfy the requirement



Amazon EFS (Linux)



Amazon FSx for Windows File Server (Windows)




Amazon  
Elastic File System  
(Amazon EFS)

Amazon EFS provides file system storage for **Linux-based** workloads.

- Fully managed elastic file system
- Scales automatically up or down as files are added and removed
- Petabytes of capacity
- Supports Network File System (NFS) protocols
  - Mount the file system to the EC2 instance
- Compatible with all Linux-based AMIs for Amazon EC2

## Common workloads and applications:

- Home directories
- File system for enterprise applications
- Application testing and development
- Database backups
- Web serving and content management
- Media workflows
- Big data analytics



Example command to mount the file system to each guest OS:

```
$ sudo mount -t nfs4 mount-target-DNS:/ ~/efs-mount-point
```

Provides fully managed shared file system storage for Microsoft Windows EC2 instances.




Amazon FSx for  
Windows File  
Server

- Native Microsoft Windows compatibility
- New Technology File System (NTFS)
- Native Server Message Block (SMB) protocol version 2.0 to 3.1.1
- Distributed File System (DFS) Namespaces and DFS Replication
- Integrates with Microsoft Active Directory and supports Windows access control lists (ACLs)
- Backed by high-performance SSD storage

Amazon FSx for Windows File Server supports a **broad set of Microsoft Windows workloads.**

- Home directories
- Lift-and-shift application workloads
- Media and entertainment workflows
- Data analytics
- Web serving and content management
- Software development environments



# Key takeaways



- Storage options for EC2 instances include
- Instance store, Amazon EBS, Amazon EFS, and Amazon FSx for Windows File Server
- For a root volume ---
- use instance store or SSD-backed Amazon EBS
- For a data volume that serves only one instance,
- Use instance store or Amazon EBS storage
- For a data volume that serves multiple Linux instances,
- use Amazon EFS
- For a data volume that serves multiple Microsoft Windows instances,
- use Amazon FSx for Windows File Server


# Amazon EC2 pricing options



# Amazon EC2 pricing options (1 of 2)

## On-Demand Instances


Pay for compute capacity by the second or by the hour with no long-term commitments.



Spiky workloads,  
workload experimentation

## Reserved Instances


Make a 1-year or 3-year commitment and receive a significant discount off on-demand prices.



Committed and  
steady-state workloads

## Savings Plans

Same discounts as Reserved Instances with more flexibility in exchange for a \$/hour commitment.




All Amazon EC2,  
AWS Fargate, and  
AWS Lambda workloads

# Amazon EC2 pricing options (2 of 2)

## Spot Instances


Spare Amazon EC2 capacity at **substantial savings** off On-Demand Instance prices.



Fault-tolerant, flexible, stateless workloads

## Dedicated Hosts

**Physical server** with Amazon EC2 instance capacity **fully dedicated for your use**.



Workloads that require the use of your own software licenses or single tenancy to meet compliance requirements

# Amazon EC2 dedicated options

Amazon EC2 dedicated options provide EC2 instance capacity on **physical servers that are dedicated for your use** (single-tenant hardware).

## Dedicated Instances


- Per-instance billing
- Automatic instance placement
- Benefit – Isolates the hosts that run your instances

## Dedicated Hosts

- Per-host billing
- Visibility of sockets, cores, and host ID
- Affinity between a host and an instance
- Targeted instance placement
- Add capacity by using an allocation request
- Benefit – Enables you to use your server-bound software licenses and address compliance requirements

# Amazon EC2 cost optimization guideline

To **optimize** the cost of Amazon EC2 instances, **combine** the available purchase options.



# Key takeaways



- Amazon EC2 pricing models include
- On-Demand Instances, Reserved Instances, Savings Plans, Spot Instances, and Dedicated Hosts
- Per-second billing is available
- For On-Demand Instances, Reserved Instances, and Spot Instances that run Amazon Linux or Ubuntu
- To optimize Amazon EC2 compute costs
- Use a combination of Reserved Instances, Savings Plans, On-Demand Instances, and Spot Instances

# Amazon EC2 cost optimization

# What is Cost Optimization?




Reduce Costs...

Pay only for *what* you need


*when* you need it.

# Lowering TCO Through Cost Optimization





# The Four Pillars of Cost Optimization



# Driver 1: Right-Sizing


## Driver 1:

Right-Sizing

Reserved Instances

Increase Elasticity

Monitor & Improve



- ❏ Select the appropriate instance types
- ❏ Downsize instances
- ❏ Leverage Amazon CloudWatch metrics

### Best practice:


- ❏ Right size, then reserve

# Optimize and Combine Amazon EC2 Purchase Types

## Driver 1:

Right-Sizing

Reserved Instances  
Increase Elasticity  
Monitor & Improve



On-Demand




Reserved



Spot



Dedicated




# Optimize and Combine Amazon EC2 Purchase Types

## Driver 1:

### Right-Sizing

Reserved Instances  
Increase Elasticity  
Monitor & Improve



### On-Demand



Spiky Workloads


### Reserved



### Spot



### Dedicated




- Pay by the hour.
- No long-term commitments

# Optimize and Combine Amazon EC2 Purchase Types

## Driver 1:

### Right-Sizing

Reserved Instances  
Increase Elasticity  
Monitor & Improve



### On-Demand



Spiky Workloads

- Pay by the hour.
- No long-term commitments

### Reserved




Steady-state Workloads

- Pay upfront
- 50-75% lower hourly rate

### Spot



### Dedicated




# Optimize and Combine Amazon EC2 Purchase Types

## Driver 1:

### Right-Sizing

Reserved Instances  
Increase Elasticity  
Monitor & Improve



### On-Demand



Spiky workloads

- Pay by the hour.
- No long-term commitments

### Reserved



Steady-state workloads

- Pay upfront
- 50-75% lower hourly rate


### Spot



Time-insensitive workloads

- Bid for unused Amazon EC2 capacity

### Dedicated




# Optimize and Combine Amazon EC2 Purchase Types

## Driver 1:

### Right-Sizing

Reserved Instances  
Increase Elasticity  
Monitor & Improve



### On-Demand



Spiky Workloads

- Pay by the hour.
- No long-term commitments

### Reserved



Steady-state Workloads

- Pay upfront
- 50-75% lower hourly rate


### Spot



Time-insensitive Workloads

- Bid for unused Amazon EC2 capacity

### Dedicated



Highly sensitive Workloads



- In your VPC
- Isolated, steady-state workloads

# Optimize and Combine Amazon EC2 Purchase Types

## Driver 1:

### Right-Sizing

Reserved Instances  
Increase Elasticity  
Monitor & Improve



- ✓ Pay only for what you use
- ✓ On-demand, elastic provisioning
- ✓ Control and security



## Driver 2:

Right-Sizing  
**Reserved Instances**  
(Ris)

Increase Elasticity  
Monitor & Improve

## Reserved Instances (RIs)/Capacity

- Amazon Elastic Compute Cloud (EC2)
- Amazon Relational Database Service (RDS)
- Amazon DynamoDB
- Amazon Redshift
- Amazon ElastiCache

## Commitment level


- 1 year
- 3 years

**Up to 75%+  
savings**

\* Dependent on specific AWS service, size/type, and region

## Driver 2:

Right-Sizing  
**Reserved Instances**  
Increase Elasticity  
Monitor & Improve



### Step 1: RI Coverage

- Cover always-on resources
- Target 70–80% always-on coverage


### Step 2: RI Utilization

- Leverage RI flexibility to increase utilization
- Merge and split RIs as needed
- Target 95% RI utilization rate

# Driver 3: Increase Elasticity

## Driver 3:

Right-Sizing  
Reserved Instances  
**Increase Elasticity**  
Monitor & Improve



### **Elasticity**

Using an instance when you need, turning it off when you don't

### **Turn off non-production instances**

Example: Dev/test

### **Auto scale production**

Use Auto Scaling to scale up and down based on demand and usage (e.g., spikes)


### **Target: 20-30% of Amazon EC2 instances**

Run in On-demand or as Spot


# Using Right-sizing and Elasticity to Lower Cost

## Driver 3:

Right-Sizing  
Reserved Instances  
**Increase Elasticity**  
Monitor & Improve



More, smaller instances vs. fewer, larger instances



\*Assumes Linux instances in the US-East (N. Virginia) Region at 720 hours per month

# Driver 4: Measure, Monitor, and Improve

## Driver 4:

Right-Sizing  
Reserved Instances  
Increase Elasticity  
Monitor & Improve



## Cost Optimization Opportunities:

1. Auto-tag resources
2. Identify always-on non production systems
3. Identify instances to downsize
4. Recommend Reserved Instance (RIs) to purchase
5. Dashboard your status
6. Consolidate your billing
7. Report on savings

# Measure, Monitor, and Improve

## Driver 4:

Right-Sizing  
Reserved Instances  
Increase Elasticity  
**Monitor & Improve**




- ❏ Optimize your AWS environment
- ❏ Reduce cost, increase performance, and improve security



- ❏ View graphs of your costs: the last 13 months
- ❏ Forecast your likely costs: the next 3 months
- ❏ View time data by day or month

# Continual Process of Cost Optimization



# AWS Lambda service



# What is AWS Lambda?




**AWS  
Lambda**

- Fully managed serverless compute
- Event-driven execution
- Sub-second metering
- Function execution limited to a maximum of 5 minutes
- Multiple languages supported

# Lambda Key Benefits



No Servers to  
Manage




Continuous  
Scaling



Sub-second  
Metering



# Getting Started with Lambda



Upload your code to  
AWS Lambda

Set up your code to trigger  
from other AWS services, HTTP  
endpoints, or in-app activity

Lambda runs your code  
only when triggered using  
only the compute  
resources needed

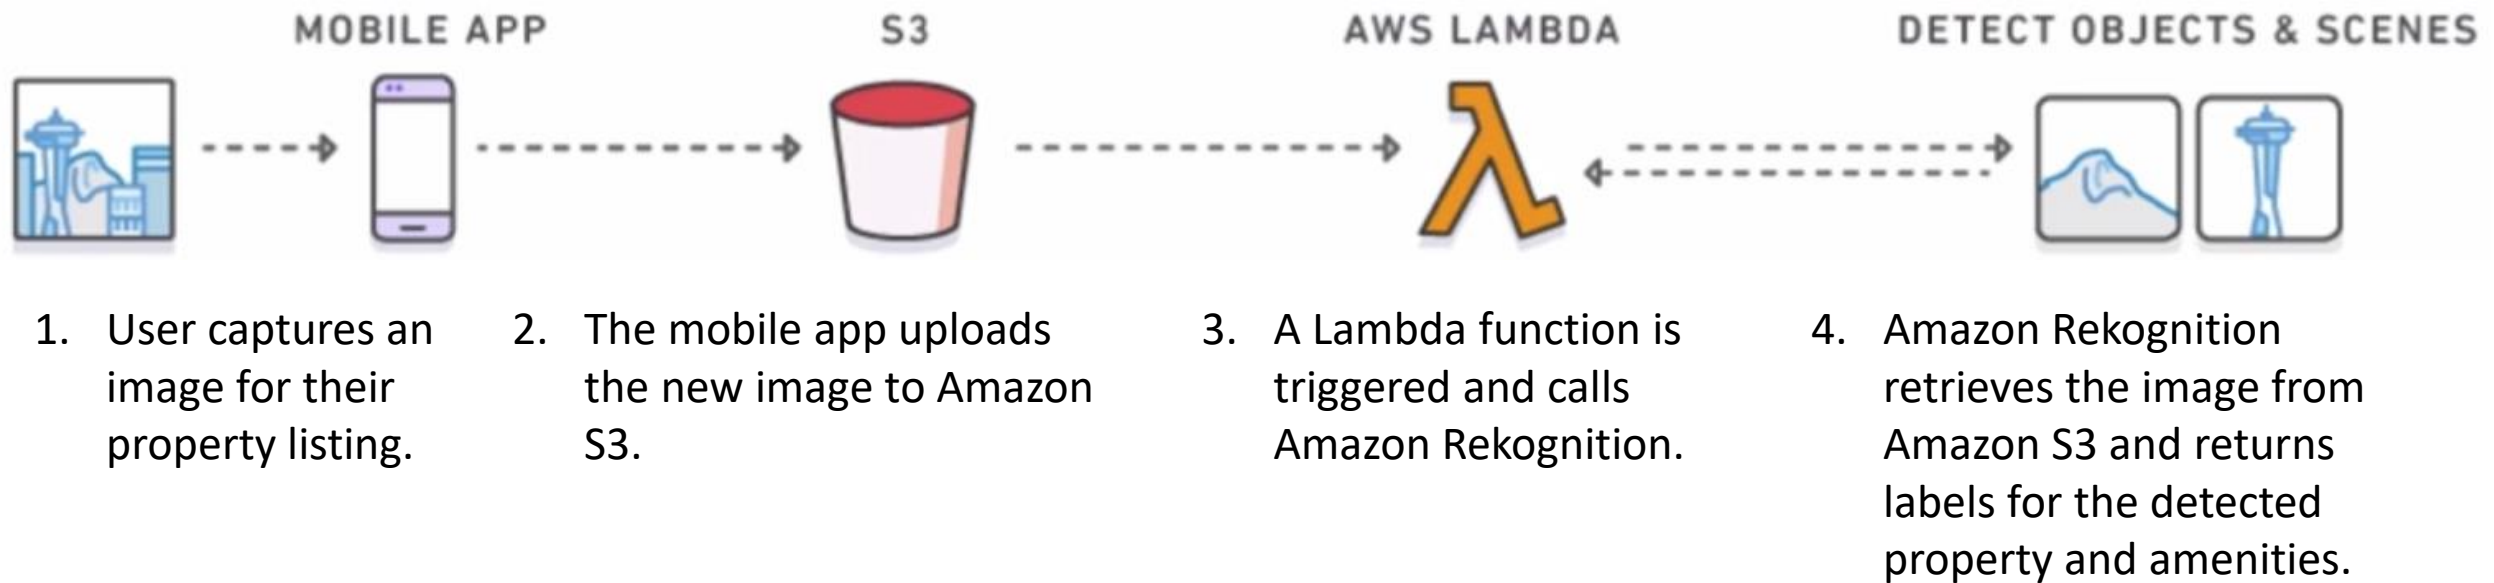
Pay just for the compute  
time you use



- Run code in response to an events.
- For example:
  - Changes to an S3 bucket
  - Changes to an Amazon Dynamo DB table
  - Respond to HTTP request
  - Invoke code with API calls
- Build serverless applications triggered by Lambda functions.
- Deploy with AWS CodePipeline and AWS CodeDeploy.



# Lambda Example




# In Review

- Fully managed serverless compute
- Event-driven execution
- Executes code only when needed and scales automatically
- Multiple languages supported



# AWS Elastic Beanstalk

# What is Elastic Beanstalk?




## **AWS Elastic Beanstalk**

- Platform as a Service (PaaS)
- Quickly deploys, scales, and manages web apps
- Reduces management complexity
- Keeps control in your hands:
  - Choose your instance type
  - Choose your database
  - Set and adjust Auto Scaling
  - Update your application
  - Access server log files
  - Enable HTTPS on load balancer



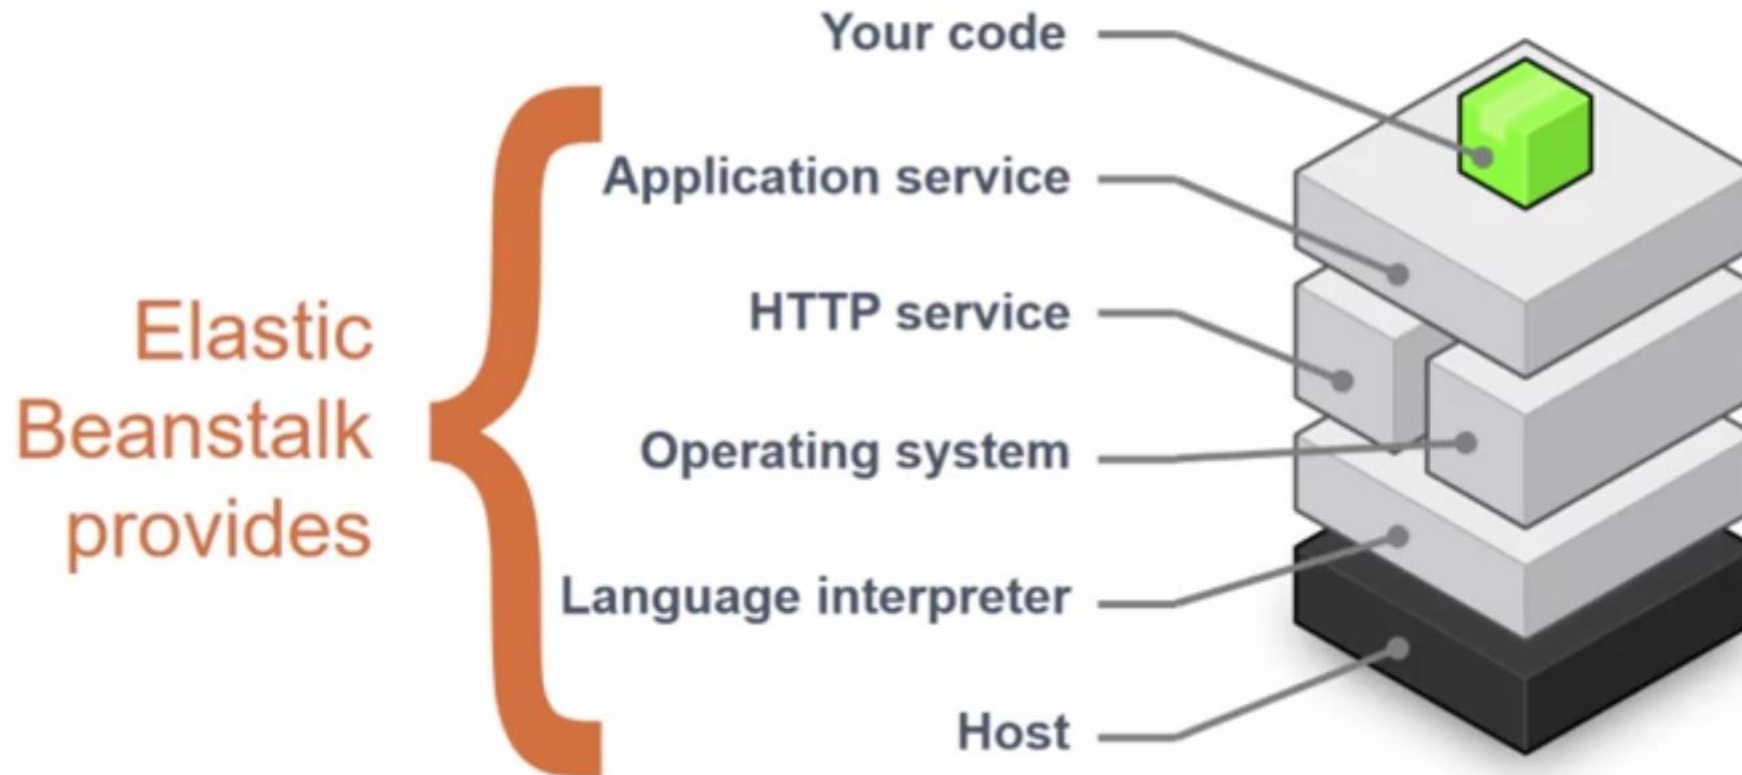
# What is Elastic Beanstalk?



**AWS**  
**Elastic**  
**Beanstalk**

- Supports a large range of platforms:
  - Packer Builder
  - Single Container, Multi-container, or Pre-configured Docker
  - Go
  - Java SE
  - Java with Tomcat
  - .NET on Windows Server with IIS
  - Node.js
  - PHP
  - Python
  - Ruby
- No charge for Elastic Beanstalk; pay only for the underlying services used.

# Elastic Beanstalk Components



# Elastic Beanstalk Key Benefits




Fast and simple to  
begin



Developer  
productivity



Impossible to  
outgrow



Complete resource  
control



- Enhances developer productivity by simplifying the process of deploying your application.
- Reduces management complexity.
- There is no charge for Elastic Beanstalk. You pay only for the services you use.



# Sample exam question

A Solutions Architect wants to design a solution to save costs for EC2 instances that do not need to run during a 2-week company shutdown. The applications running on the instances store data in instance memory (RAM) that must be present when the instances resume operation.

Which approach should the Solutions Architect recommend to shut down and resume the instances?

- A. Modify the application to store the data on instance store volumes. Reattach the volumes while restarting them.
- B. Snapshot the instances before stopping them. Restore the snapshot after restarting the instances.
- C. Run the applications on instances enabled for hibernation. Hibernate the instances before the shutdown.
- D. Note the Availability Zone for each instance before stopping it. Restart the instances in the same Availability Zones after the shutdown.

# Additional resources

- [Amazon EC2 User Guide for Linux Instances](#)
- [Amazon EC2 User Guide for Windows Instances](#)
- [Amazon EC2 FAQs](#)
- [EC2 Image Builder User Guide](#)
- [EC2 Image Builder FAQs](#)
- [AWS Compute Optimizer User Guide](#)
- [AWS Compute Optimizer FAQs](#)
- [How AWS Pricing Works](#)